


300 prompts · 10 countries · 2 VLM judges × 3 presentation orders

kept gain **-0.060**

A useful gate should keep prompts where the judge does better than on the prompts it rejects. Unanimity does (+0.141 vs. -0.030); cross-model agreement does the reverse (-0.060 vs. +0.064).

[illegible]