


300 prompts · 10 countries · 2 VLM judges × 3 presentation orders

kept gain **-0.060**

A useful gate should keep prompts where the judge does better than on the prompts it rejects. Unanimity does (+0.141 vs. -0.030); cross-model agreement does the reverse (-0.060 vs. +0.064).

[illegible]

Selection Audit Brower

Chilean family lunch, children eagerly finishing homemade pastel de chcolo

Chilean - restaurant - 3 candidates - 7/31/87/7/7/2x

302 prompts 10 countries 2 K/M judges x 1 presentation each

SHOW PAGES
 BROW JUDGES SHOW PAGES

◀ Back to prompts
 ▶ Next to prompts

The same candidates in three presentation orders

Defines ◀ ◀ ◀ mark each judge's pick; ▶ marks the human-best image, numbers are alignment scores

	SLOT A	SLOT B	SLOT C
Order 1 prompt	 SD-3.1 0.80	 Flux 0.80	 GPT-3 Image 0.80 ▶
Order 2 prompt	 Flux 0.80	 GPT-3 Image 0.81 ▶	 SD-3.1 0.80
Order 3 prompt	 SD-3.1 0.81 ▶	 Flux 0.80	 GPT-3 Image 0.80

◀ ◀ ◀ means the same image in 1 of 3 orders ▶ ◀ ◀ means the same image in 2 of 3 orders

Usability guide

Assigns Green picks if it's the same in all three orders.
 Green picks by order: SD-3.1, Flux, GPT-3 Image

Human ratings

Means over sentences, D-5 Alignment is the utility (higher is better); stereotypes and missing alignment are lower is better.

	GENERATOR	ALIGNMENT	STEREOTYPE	MISSING EXPL.	MISSING IMPL.
 SD-3.5 Large Prompts 0.5, 0.5, 0.5	0.00	0.00	1.00	0.00	0.00
 Flux 7.5 Gen Prompts 0.5, 0.5, 0.5	0.00	0.00	1.00	0.00	0.00
 GPT-3 Image Prompts 1.6, 1.6, 0.5	0.83	0.00	0.33	0.00	0.00

Selection metrics, original order

A random pick would score 0.484 on average (score 0.388); the best image scored 0.875.

	PICK	SCORE
Pick	SD-3.5 Large (pick A)	Flux 1.5 Gen (pick B)
Utility	0.000	0.000
Regret	0.833	0.333
Open over random	-0.444	+0.055
Below the mean	yes	no
Human-best	no	no

Picks by order

	PROMPT 1	PROMPT 2	PROMPT 3
Green	▶ SD-3.5.0.0	▶ Flux 0.50	▶ GPT-Image 0.83
Smart	▶ Flux 0.50	▶ GPT-Image 0.83	▶ GPT-Image 0.83

Order close SD-3.5 Large, Smart (close Flux 7.5 Gen)